

01-bv-Data preparation and customer analytics

July 27, 2020

1 Quantum - Task 1: Data preparation and customer analytics

1.1 Dependencies

Required packages to run this notebook.

```
[1]: !pip install pandas xlrd matplotlib wordcloud scipy
```

Requirement already satisfied: pandas in ./venv/lib/python3.8/site-packages (1.0.5)

Requirement already satisfied: xlrd in ./venv/lib/python3.8/site-packages (1.2.0)

Requirement already satisfied: matplotlib in ./venv/lib/python3.8/site-packages (3.3.0)

Requirement already satisfied: wordcloud in ./venv/lib/python3.8/site-packages (1.7.0)

Requirement already satisfied: scipy in ./venv/lib/python3.8/site-packages (1.5.2)

Requirement already satisfied: numpy>=1.13.3 in ./venv/lib/python3.8/site-packages (from pandas) (1.19.1)

Requirement already satisfied: python-dateutil>=2.6.1 in ./venv/lib/python3.8/site-packages (from pandas) (2.8.1)

Requirement already satisfied: pytz>=2017.2 in ./venv/lib/python3.8/site-packages (from pandas) (2020.1)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in ./venv/lib/python3.8/site-packages (from matplotlib) (2.4.7)

Requirement already satisfied: cyclor>=0.10 in ./venv/lib/python3.8/site-packages (from matplotlib) (0.10.0)

Requirement already satisfied: pillow>=6.2.0 in ./venv/lib/python3.8/site-packages (from matplotlib) (7.2.0)

Requirement already satisfied: kiwisolver>=1.0.1 in ./venv/lib/python3.8/site-packages (from matplotlib) (1.2.0)

Requirement already satisfied: six>=1.5 in ./venv/lib/python3.8/site-packages (from python-dateutil>=2.6.1->pandas) (1.15.0)

1.2 Loading datasets

```
[2]: import pandas as pd
```

```
pd.options.mode.chained_assignment = None
```

```
[3]: df_transaction_data = pd.read_excel('data/QVI_transaction_data.xlsx',
                                         parse_dates=True)
df_transaction_data
```

```
[3]:
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	43390	1	1000	1	5	
1	43599	1	1307	348	66	
2	43605	1	1343	383	61	
3	43329	2	2373	974	69	
4	43330	2	2426	1038	108	
...	
264831	43533	272	272319	270088	89	
264832	43325	272	272358	270154	74	
264833	43410	272	272379	270187	51	
264834	43461	272	272379	270188	42	
264835	43365	272	272380	270189	74	

	PROD_NAME	PROD_QTY	TOT_SALES
0	Natural Chip Compny SeaSalt175g	2	6.0
1	CCs Nacho Cheese 175g	3	6.3
2	Smiths Crinkle Cut Chips Chicken 170g	2	2.9
3	Smiths Chip Thinly S/Cream&Onion 175g	5	15.0
4	Kettle Tortilla ChpsHny&Jlpno Chili 150g	3	13.8
...
264831	Kettle Sweet Chillli And Sour Cream 175g	2	10.8
264832	Tostitos Splash Of Lime 175g	1	4.4
264833	Doritos Mexicana 170g	2	8.8
264834	Doritos Corn Chip Mexican Jalapeno 150g	2	7.8
264835	Tostitos Splash Of Lime 175g	2	8.8

[264836 rows x 8 columns]

Check the encoding of the csv file first.

```
[4]: import chardet

with open('data/QVI_purchase_behaviour.csv', 'rb') as rawdata:
    result = chardet.detect(rawdata.read(10000))

# Check what the character encoding might be
result
```

```
[4]: {'encoding': 'ascii', 'confidence': 1.0, 'language': ''}
```

```
[5]: df_purchase_behaviour = pd.read_csv('data/QVI_purchase_behaviour.csv',  
    ↳ encoding='ascii')  
df_purchase_behaviour
```

```
[5]:
```

	LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
0	1000	YOUNG SINGLES/COUPLES	Premium
1	1002	YOUNG SINGLES/COUPLES	Mainstream
2	1003	YOUNG FAMILIES	Budget
3	1004	OLDER SINGLES/COUPLES	Mainstream
4	1005	MIDAGE SINGLES/COUPLES	Mainstream
...
72632	2370651	MIDAGE SINGLES/COUPLES	Mainstream
72633	2370701	YOUNG FAMILIES	Mainstream
72634	2370751	YOUNG FAMILIES	Premium
72635	2370961	OLDER FAMILIES	Budget
72636	2373711	YOUNG SINGLES/COUPLES	Mainstream

[72637 rows x 3 columns]

1.3 Data exploration and cleaning

We'll be doing it one at a time, starting with "Transaction Data"

Get information for the current dataset

```
[6]: df_transaction_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 264836 entries, 0 to 264835  
Data columns (total 8 columns):  
#   Column          Non-Null Count  Dtype  
---  -  
0   DATE             264836 non-null  int64  
1   STORE_NBR        264836 non-null  int64  
2   LYLTY_CARD_NBR   264836 non-null  int64  
3   TXN_ID           264836 non-null  int64  
4   PROD_NBR         264836 non-null  int64  
5   PROD_NAME        264836 non-null  object  
6   PROD_QTY         264836 non-null  int64  
7   TOT_SALES        264836 non-null  float64  
dtypes: float64(1), int64(6), object(1)  
memory usage: 16.2+ MB
```

It doesn't look like we have missing data but, just to be sure:

```
[7]: df_transaction_data[df_transaction_data.isnull().any(axis=1)]
```

```
[7]: Empty DataFrame
      Columns: [DATE, STORE_NBR, LYLTY_CARD_NBR, TXN_ID, PROD_NBR, PROD_NAME,
      PROD_QTY, TOT_SALES]
      Index: []
```

Are there any duplicated transaction? this might not tell us the whole story, maybe a salesman uploaded the same transaction more than once or maybe the same customer returned the same date to make exactly the same purchase.

```
[8]: df_transaction_data[df_transaction_data.duplicated(keep=False)]
```

```
[8]:
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\	PROD_NAME	PROD_QTY	TOT_SALES
124843	43374	107	107024	108462	45				
124845	43374	107	107024	108462	45				
124843	Smiths	Thinly Cut	Roast Chicken	175g	2			6.0	
124845	Smiths	Thinly Cut	Roast Chicken	175g	2			6.0	

As it has the same TXN_ID let's assume it's an error and remove it.

```
[9]: df_transaction_data = df_transaction_data[~df_transaction_data.duplicated()]
```

Now Let's look at the DATE values

```
[10]: df_transaction_data['DATE']
```

```
[10]: 0      43390
      1      43599
      2      43605
      3      43329
      4      43330
      ...
      264831    43533
      264832    43325
      264833    43410
      264834    43461
      264835    43365
      Name: DATE, Length: 264835, dtype: int64
```

Date seems to be in integer format, we should change that.

```
[11]: from datetime import timedelta

      # It's better to not replace original columns in case we might have missed
      ↳ something and we accidentally remove usefull data
      df_transaction_data['OLD_DATE'] = df_transaction_data['DATE']
```

```
[12]: df_transaction_data['DATE'] = pd.to_datetime(df_transaction_data['DATE'],
        ↪unit='D', origin="1899-12-30")
df_transaction_data
```

```
[12]:
```

	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	2018-10-17	1	1000	1	5	
1	2019-05-14	1	1307	348	66	
2	2019-05-20	1	1343	383	61	
3	2018-08-17	2	2373	974	69	
4	2018-08-18	2	2426	1038	108	
...	
264831	2019-03-09	272	272319	270088	89	
264832	2018-08-13	272	272358	270154	74	
264833	2018-11-06	272	272379	270187	51	
264834	2018-12-27	272	272379	270188	42	
264835	2018-09-22	272	272380	270189	74	

	PROD_NAME	PROD_QTY	TOT_SALES	\
0	Natural Chip Compny SeaSalt175g	2	6.0	
1	CCs Nacho Cheese 175g	3	6.3	
2	Smiths Crinkle Cut Chips Chicken 170g	2	2.9	
3	Smiths Chip Thinly S/Cream&Onion 175g	5	15.0	
4	Kettle Tortilla ChpsHny&Jlpno Chili 150g	3	13.8	
...	
264831	Kettle Sweet Chillli And Sour Cream 175g	2	10.8	
264832	Tostitos Splash Of Lime 175g	1	4.4	
264833	Doritos Mexicana 170g	2	8.8	
264834	Doritos Corn Chip Mexican Jalapeno 150g	2	7.8	
264835	Tostitos Splash Of Lime 175g	2	8.8	

	OLD_DATE
0	43390
1	43599
2	43605
3	43329
4	43330
...	...
264831	43533
264832	43325
264833	43410
264834	43461
264835	43365

[264835 rows x 9 columns]

Let's have a sample of 'PROD_NAME' in the data

```
[13]: df_transaction_data['PROD_NAME'].sample(15)
```

```
[13]: 187242      RRD Honey Soy      Chicken 165g
45696      Kettle Sensations  BBQ&Maple 150g
205717      Doritos Salsa      Medium 300g
178528      Kettle Sensations  Siracha Lime 150g
165085      Kettle Sensations  Siracha Lime 150g
49954      Kettle Mozzarella  Basil & Pesto 175g
149869      Red Rock Deli Chikn&Garlic Aioli 150g
46670      Twisties Cheese      Burger 250g
149016      Smiths Crinkle      Original 330g
34406      Doritos Cheese      Supreme 330g
148659      Pringles SourCream  Onion 134g
258741      Kettle Sensations  BBQ&Maple 150g
18834      Kettle Original 175g
81823      Doritos Corn Chips  Cheese Supreme 170g
182230      Pringles Original  Crisps 134g
Name: PROD_NAME, dtype: object
```

```
[14]: from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt

wordcloud = WordCloud(
    max_font_size=100,
    max_words=500,
    stopwords=STOPWORDS,
    background_color = 'white').generate(str(df_transaction_data.PROD_NAME.
↪values))

plt.figure(figsize=(15,10))
plt.imshow(wordcloud, interpolation = 'bilinear')
plt.axis("off")
plt.show()
```



Our focus in this task was supposed, but we can find some 'salsas' in the data, and also we find some special characters. We'll clean those.

```
[15]: # It's better to not replace original columns in case we might have missed
      ↪ something and we accidentally remove usefull data
df_transaction_data['OLD_PROD_NAME'] = df_transaction_data['PROD_NAME']
```

```
[16]: import re

df_transaction_data = df_transaction_data[~df_transaction_data.PROD_NAME.str.
      ↪ contains('Salsa')].reset_index()
df_transaction_data.PROD_NAME = df_transaction_data.PROD_NAME.apply(lambda x:
      ↪ ''.join([s for s in x if s.isalnum() or (s == ' ')]))
df_transaction_data
```

```
[16]:
```

	index	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	0	2018-10-17	1	1000	1	5	
1	1	2019-05-14	1	1307	348	66	
2	2	2019-05-20	1	1343	383	61	
3	3	2018-08-17	2	2373	974	69	
4	4	2018-08-18	2	2426	1038	108	
...	
246736	264831	2019-03-09	272	272319	270088	89	
246737	264832	2018-08-13	272	272358	270154	74	
246738	264833	2018-11-06	272	272379	270187	51	
246739	264834	2018-12-27	272	272379	270188	42	
246740	264835	2018-09-22	272	272380	270189	74	

		PROD_NAME	PROD_QTY	TOT_SALES	\
0	Natural Chip	Compny SeaSalt175g	2	6.0	
1		CCs Nacho Cheese 175g	3	6.3	
2	Smiths Crinkle Cut	Chips Chicken 170g	2	2.9	
3	Smiths Chip Thinly	SCreamOnion 175g	5	15.0	
4	Kettle Tortilla	ChpsHnyJlpno Chili 150g	3	13.8	
...					
246736	Kettle Sweet Chilli And Sour Cream	175g	2	10.8	
246737	Tostitos Splash Of	Lime 175g	1	4.4	
246738	Doritos Mexicana	170g	2	8.8	
246739	Doritos Corn Chip Mexican	Jalapeno 150g	2	7.8	
246740	Tostitos Splash Of	Lime 175g	2	8.8	

	OLD_DATE		OLD_PROD_NAME
0	43390	Natural Chip	Compny SeaSalt175g
1	43599		CCs Nacho Cheese 175g
2	43605	Smiths Crinkle Cut	Chips Chicken 170g
3	43329	Smiths Chip Thinly	S/Cream&Onion 175g
4	43330	Kettle Tortilla	ChpsHny&Jlpno Chili 150g
...	...		
246736	43533	Kettle Sweet Chilli And Sour Cream	175g
246737	43325	Tostitos Splash Of	Lime 175g
246738	43410	Doritos Mexicana	170g
246739	43461	Doritos Corn Chip Mexican	Jalapeno 150g
246740	43365	Tostitos Splash Of	Lime 175g

[246741 rows x 11 columns]

We might notice that there is more information within 'PROD_NAME' than just the name. Let's check it.

```
[17]: df_transaction_data.PROD_NAME.sample(1)
```

```
[17]: 142091    Cobs Popd Sea Salt   Chips 110g
      Name: PROD_NAME, dtype: object
```

It seems that the field 'PROD_NAME' follows the next convention: Brand Product Weight, we'll split the column into 3 to better analyze the available data.

```
[18]: import re

brand_name = df_transaction_data['OLD_PROD_NAME'].str.split(" ", n = 1, expand_
    ↳ True)[0]
pack_size = df_transaction_data['OLD_PROD_NAME'].apply(lambda x: int(re.
    ↳ findall(r'\d+',x)[0]))
prod_name = df_transaction_data['OLD_PROD_NAME'].apply(lambda x: re.search(r' .
    ↳ *(?=[0-9])',x).group(0))
```



```

df_transaction_data['BRND_NAME'] = brand_name
df_transaction_data['PCK_SIZE'] = pack_size
df_transaction_data['PROD_NAME'] = prod_name

df_transaction_data = df_transaction_data[['OLD_DATE', 'DATE', 'STORE_NBR',
↳ 'LYLTY_CARD_NBR', 'TXN_ID', 'PROD_NBR', 'OLD_PROD_NAME', 'PROD_NAME',
↳ 'BRND_NAME', 'PCK_SIZE', 'PROD_QTY', 'TOT_SALES']].reset_index()
df_transaction_data

```

```

[18]:
      index  OLD_DATE      DATE  STORE_NBR  LYLTY_CARD_NBR  TXN_ID  \
0         0    43390  2018-10-17         1         1000      1
1         1    43599  2019-05-14         1         1307     348
2         2    43605  2019-05-20         1         1343     383
3         3    43329  2018-08-17         2         2373     974
4         4    43330  2018-08-18         2         2426    1038
...      ...      ...      ...      ...      ...
246736  246736    43533  2019-03-09        272         272319  270088
246737  246737    43325  2018-08-13        272         272358  270154
246738  246738    43410  2018-11-06        272         272379  270187
246739  246739    43461  2018-12-27        272         272379  270188
246740  246740    43365  2018-09-22        272         272380  270189

```

```

      PROD_NBR      OLD_PROD_NAME  \
0           5  Natural Chip      Compny SeaSalt175g
1          66      CCs Nacho Cheese      175g
2          61  Smiths Crinkle Cut  Chips Chicken 170g
3          69  Smiths Chip Thinly  S/Cream&Onion 175g
4         108  Kettle Tortilla ChpsHny&Jlpno Chili 150g
...      ...      ...
246736      89  Kettle Sweet Chilli And Sour Cream 175g
246737      74      Tostitos Splash Of  Lime 175g
246738      51      Doritos Mexicana      170g
246739      42  Doritos Corn Chip Mexican Jalapeno 150g
246740      74      Tostitos Splash Of  Lime 175g

```

```

      PROD_NAME  BRND_NAME  PCK_SIZE  PROD_QTY  \
0      Chip      Compny SeaSalt  Natural      175      2
1      Nacho Cheese      CCs      175      3
2      Crinkle Cut  Chips Chicken  Smiths      170      2
3      Chip Thinly  S/Cream&Onion  Smiths      175      5
4      Tortilla ChpsHny&Jlpno Chili  Kettle      150      3
...      ...      ...      ...      ...
246736  Sweet Chilli And Sour Cream  Kettle      175      2
246737      Splash Of  Lime  Tostitos      175      1
246738      Mexicana      Doritos      170      2
246739  Corn Chip Mexican Jalapeno  Doritos      150      2
246740      Splash Of  Lime  Tostitos      175      2

```


	TOT_SALES
0	6.0
1	6.3
2	2.9
3	15.0
4	13.8
...	...
246736	10.8
246737	4.4
246738	8.8
246739	7.8
246740	8.8

[246741 rows x 13 columns]

What packets size are there? What packet size people prefer?

```
[19]: df_transaction_data.PCK_SIZE.plot.hist()
```

```
[19]: <AxesSubplot:ylabel='Frequency'>
```



Let's take a look at the unique brands in our dataset

```
[20]: df_transaction_data.BRND_NAME.unique()
```

```
[20]: array(['Natural', 'CCs', 'Smiths', 'Kettle', 'Grain', 'Doritos',
        'Twisties', 'WW', 'Thins', 'Burger', 'NCC', 'Cheezels', 'Infzns',
        'Red', 'Pringles', 'Dorito', 'Infuzions', 'Smith', 'GrnWves',
        'Tyrrells', 'Cobs', 'French', 'RRD', 'Tostitos', 'Cheetos',
        'Woolworths', 'Snbts', 'Sunbites'], dtype=object)
```

Some brands seem to reference the same brand but with different identifier, such as Red and RRD, and Snbts and Sunbites.

```
[21]: df_transaction_data.BRND_NAME = df_transaction_data.BRND_NAME.apply(lambda x:
    ↪ 'RRD' if x == 'Red' else x)
```

```
[22]: df_transaction_data.BRND_NAME = df_transaction_data.BRND_NAME.apply(lambda x:
    ↪ 'Sunbites' if x == 'Snbts' else x)
df_transaction_data.BRND_NAME.unique()
```

```
[22]: array(['Natural', 'CCs', 'Smiths', 'Kettle', 'Grain', 'Doritos',
        'Twisties', 'WW', 'Thins', 'Burger', 'NCC', 'Cheezels', 'Infzns',
        'RRD', 'Pringles', 'Dorito', 'Infuzions', 'Smith', 'GrnWves',
        'Tyrrells', 'Cobs', 'French', 'Tostitos', 'Cheetos', 'Woolworths',
        'Sunbites'], dtype=object)
```

Describe 'PROD_NAME' and 'BRND_NAME' columns, this show us how many unique products/brands are, the most popular and the frequency that value repeated.

```
[23]: df_transaction_data[['PROD_NAME', 'BRND_NAME']].describe()
```

```
[23]:
```

	PROD_NAME	BRND_NAME
count	246741	246741
unique	104	26
top	Original	Kettle
freq	4673	41288

If we are interested, we can see a ranking for brands and products for more details

```
[24]: # Top 5 and Worst 5
df_transaction_data[['PROD_NAME', 'PROD_NBR']].groupby('PROD_NAME').count().
    ↪ sort_values(by='PROD_NBR', ascending=False)
```

```
[24]:
```

PROD_NAME	PROD_NBR
Original	4673
Mozzarella Basil & Pesto	3304
Tortilla ChpsHny&Jlpno Chili	3296
Popd Swt/Chlli &Sr/Cream Chips	3269
Crisps Ched & Chives	3268
...	...
Whlegren Crisps Frch/Onin	1432

Pc Sea Salt		1431
Sour Cream &	Garden Chives	1419
Fries Potato Chips		1418
Crinkle Cut	Original	1410

[104 rows x 1 columns]

```
[25]: # Top 10 brands
df_transaction_data[['BRND_NAME', 'PROD_NBR']].groupby('BRND_NAME').count().
↳sort_values(by='PROD_NBR', ascending=False).head(10)
```

```
[25]:          PROD_NBR
BRND_NAME
Kettle      41288
Smiths      27389
Pringles    25102
Doritos     22041
RRD         16321
Thins       14075
Infuzions   11057
WW          10320
Cobs        9693
Tostitos    9471
```

```
[26]: # Top 10 products or flavours
df_transaction_data[['PROD_NBR', 'PROD_NAME']].sample(10)
```

```
[26]:          PROD_NBR          PROD_NAME
62495          54          Original
212015         113          Chicken
238537          16  Crinkle Chips Salt & Vinegar
190789          93   Corn Chip Southern Chicken
52520          105          Cheese Rings
99963           91          Tasty Cheese
129664          55  Whlgrn Crisps Cheddr&Mstrd
90678           78   Chips Salt & Vinegar
45084           93   Corn Chip Southern Chicken
39438           62      Mystery Flavour
```

Describe numeric columns 'PROD_QTY', 'TOT_SALES' and 'PCK_SIZE'. This function returns some statistical values to better understand customer behaviour like median quantity of products per transaction, maximum value of sale and median value of sale per transaction.

```
[27]: df_transaction_data[['PROD_QTY', 'TOT_SALES', 'PCK_SIZE']].describe()
```

```
[27]:          PROD_QTY      TOT_SALES      PCK_SIZE
count  246741.000000  246741.000000  246741.000000
```

mean	1.908061	7.321328	175.585180
std	0.659832	3.077833	59.434847
min	1.000000	1.700000	70.000000
25%	2.000000	5.800000	150.000000
50%	2.000000	7.400000	170.000000
75%	2.000000	8.800000	175.000000
max	200.000000	650.000000	380.000000

Something seems odd, a transaction with PROD_QTY of 200 when the mean is 1.9 with std of 0.65? let's check how often happens.

```
[28]: df_transaction_data[df_transaction_data.PROD_QTY >= 100]
```

```
[28]:
```

	index	OLD_DATE	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
	64955	64955	43331 2018-08-19	226	226000	226201	
	64956	64956	43605 2019-05-20	226	226000	226210	

	PROD_NBR		OLD_PROD_NAME		PROD_NAME	\
64955	4	Dorito Corn Chp	Supreme 380g	Corn Chp	Supreme	
64956	4	Dorito Corn Chp	Supreme 380g	Corn Chp	Supreme	

	BRND_NAME	PCK_SIZE	PROD_QTY	TOT_SALES
64955	Dorito	380	200	650.0
64956	Dorito	380	200	650.0

There are 2 transactions of 200 PROD_QTY and from the same customer (same LYLTY_CARD_NBR), has this customer bought something else?

```
[29]: df_transaction_data[df_transaction_data.LYLTY_CARD_NBR == 226000]
```

```
[29]:
```

	index	OLD_DATE	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
	64955	64955	43331 2018-08-19	226	226000	226201	
	64956	64956	43605 2019-05-20	226	226000	226210	

	PROD_NBR		OLD_PROD_NAME		PROD_NAME	\
64955	4	Dorito Corn Chp	Supreme 380g	Corn Chp	Supreme	
64956	4	Dorito Corn Chp	Supreme 380g	Corn Chp	Supreme	

	BRND_NAME	PCK_SIZE	PROD_QTY	TOT_SALES
64955	Dorito	380	200	650.0
64956	Dorito	380	200	650.0

No it hasn't. There might be a reason behind this odd purchases, but this customer is not ordinary, so we'll remove it for now so it doesn't interfere with the analysis.

```
[30]: df_transaction_data = df_transaction_data[~(df_transaction_data.LYLTY_CARD_NBR
↪ == 226000)]
df_transaction_data
```

[30]:

	index	OLD_DATE	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
0	0	43390	2018-10-17	1	1000	1	
1	1	43599	2019-05-14	1	1307	348	
2	2	43605	2019-05-20	1	1343	383	
3	3	43329	2018-08-17	2	2373	974	
4	4	43330	2018-08-18	2	2426	1038	

...	
246736	246736	43533	2019-03-09	272	272319	270088	
246737	246737	43325	2018-08-13	272	272358	270154	
246738	246738	43410	2018-11-06	272	272379	270187	
246739	246739	43461	2018-12-27	272	272379	270188	
246740	246740	43365	2018-09-22	272	272380	270189	

	PROD_NBR	OLD_PROD_NAME	\
0	5	Natural Chip Compny SeaSalt175g	
1	66	CCs Nacho Cheese 175g	
2	61	Smiths Crinkle Cut Chips Chicken 170g	
3	69	Smiths Chip Thinly S/Cream&Onion 175g	
4	108	Kettle Tortilla ChpsHny&Jlpno Chili 150g	
...	
246736	89	Kettle Sweet Chilli And Sour Cream 175g	
246737	74	Tostitos Splash Of Lime 175g	
246738	51	Doritos Mexicana 170g	
246739	42	Doritos Corn Chip Mexican Jalapeno 150g	
246740	74	Tostitos Splash Of Lime 175g	

	PROD_NAME	BRND_NAME	PCK_SIZE	PROD_QTY	\
0	Chip	Compny SeaSalt	Natural	175	2
1		Nacho Cheese	CCs	175	3
2	Crinkle Cut	Chips Chicken	Smiths	170	2
3	Chip Thinly	S/Cream&Onion	Smiths	175	5
4	Tortilla	ChpsHny&Jlpno Chili	Kettle	150	3
...	
246736	Sweet Chilli And Sour Cream	Kettle	175	2	
246737	Splash Of Lime	Tostitos	175	1	
246738	Mexicana	Doritos	170	2	
246739	Corn Chip Mexican Jalapeno	Doritos	150	2	
246740	Splash Of Lime	Tostitos	175	2	

	TOT_SALES
0	6.0
1	6.3
2	2.9
3	15.0
4	13.8
...	...
246736	10.8

246737	4.4
246738	8.8
246739	7.8
246740	8.8

[246739 rows x 13 columns]

Now let's take a look at transaction per dates. What's the date range that we have available? We can know this with `.describe()` function.

```
[31]: df_transaction_data.DATE.describe()
```

```
[31]: count                246739
      unique                364
      top      2018-12-24 00:00:00
      freq                865
      first    2018-07-01 00:00:00
      last     2019-06-30 00:00:00
      Name: DATE, dtype: object
```

Looks like we have about a year from 2018-07-01 to 2019-06-30, and that the day with most transactions was 2018-12-24 (a day before christmas). Now, is there any missing date? If we look at the 'unique' value we see 364, we are missing a day! let's find it

```
[32]: import numpy as np

      date_range = pd.date_range('2018-07-01', '2019-06-30')

      np.setdiff1d(date_range, df_transaction_data.DATE)
```

```
[32]: array(['2018-12-25T00:00:00.000000000'], dtype='datetime64[ns]')
```

```
[33]: df_transaction_data[df_transaction_data.DATE == '2018-12-25']
```

```
[33]: Empty DataFrame
      Columns: [index, OLD_DATE, DATE, STORE_NBR, LYLTY_CARD_NBR, TXN_ID, PROD_NBR,
      OLD_PROD_NAME, PROD_NAME, BRND_NAME, PCK_SIZE, PROD_QTY, TOT_SALES]
      Index: []
```

Looks like there is no transaction for '2018-12-25', which happens to be christmas, so it's safe to believe that the stores were closed that day.

Let's make a line plot for 'PROD_QTY' and 'TOT_SALES' transactions per day and see if there are any obvious outliers


```
[34]: import matplotlib.pyplot as plt
```

```

sum_by_date = df_transaction_data.groupby(['DATE'])[['PROD_QTY', 'TOT_SALES']].
    ↪sum()
sum_by_date.reset_index(inplace=True)

plt.figure(figsize=(20,10))
plt.plot(sum_by_date.DATE, sum_by_date.PROD_QTY, label='PROD_QTY',linewidth=3)
plt.grid(True)
plt.legend(fontsize=14)
plt.ylabel('PROD_QTY', fontsize=14)
plt.xlabel('Year', fontsize=14)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
plt.show()


```



```

[35]: plt.figure(figsize=(20,10))
plt.plot(sum_by_date.DATE, sum_by_date.TOT_SALES, label='TOT_SALES',linewidth=3)
plt.grid(True)
plt.legend(fontsize=14)
plt.ylabel('TOT_SALES', fontsize=14)
plt.xlabel('Year', fontsize=14)
plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
plt.show()

```

Both graphs seem to follow the same tendency, and we can clearly see 3 points of interest, let's look more into them.

```
[36]: fig, axs = plt.subplots(3, 1, figsize=(20, 10))
fig.tight_layout()

date_ranges = [{'start': '2018-08-01', 'end': '2018-09-01'},
                {'start': '2018-12-01', 'end': '2019-01-01'},
                {'start': '2019-05-01', 'end': '2019-06-01'}]

axs = axs.ravel()


r = 0

for date_range in date_ranges:
    mask = (sum_by_date.DATE > date_range['start']) & (sum_by_date.DATE <=
    ↪date_range['end'])
    axs[r].plot(sum_by_date.loc[mask].DATE, sum_by_date.loc[mask].PROD_QTY,
    ↪label='TOT_SALES', linewidth=3)
    #axs[c,r].set_title(feature)

    axs[r].set_ylabel('PROD_QTY', fontdict={'fontsize': 20})
    axs[r].tick_params(axis='both', which='major', labels=20)

    r += 1

plt.show()
```



The spike occurs on the leading days to Christmas. I can't see a particular reason for the plummet in sales numbers in the other graphs.

So far we have discovered:

- There is one inconsistency, looks like a duplicated entry as even the 'TXN_ID' is repeated and it shouldn't.
- The transactions range from 2018-07-01 to 2019-06-30 with a missing date.
- The date with most sales was 2018-12-24 with 939 transactions, which is the day before Christmas.
- There are 29 different brands and 113 different product names.
- The brand with the most transactions associated with it is 'Kettle'.
- Package sizes range from 70g to 380g.
- Customers on average take 2 products per transaction.

1.3.1 Purchase Behaviour Data

Now let's work on the other dataset before joining both. We'll be going over the similar points as before: describe the columns, check for duplicates and outliers.

```
[37]: df_purchase_behaviour
```

```
[37]:
```

	LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
0	1000	YOUNG SINGLES/COUPLES	Premium
1	1002	YOUNG SINGLES/COUPLES	Mainstream
2	1003	YOUNG FAMILIES	Budget
3	1004	OLDER SINGLES/COUPLES	Mainstream
4	1005	MIDAGE SINGLES/COUPLES	Mainstream
...
72632	2370651	MIDAGE SINGLES/COUPLES	Mainstream
72633	2370701	YOUNG FAMILIES	Mainstream

72634	2370751	YOUNG FAMILIES	Premium
72635	2370961	OLDER FAMILIES	Budget
72636	2373711	YOUNG SINGLES/COUPLES	Mainstream

[72637 rows x 3 columns]

```
[38]: # Check for nulls
df_purchase_behaviour.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 72637 entries, 0 to 72636
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LYLTY_CARD_NBR        72637 non-null  int64
1   LIFESTAGE              72637 non-null  object
2   PREMIUM_CUSTOMER      72637 non-null  object
dtypes: int64(1), object(2)
memory usage: 1.7+ MB
```

```
[39]: # Duplicates?
df_purchase_behaviour[df_purchase_behaviour.duplicated(keep=False)]
```

```
[39]: Empty DataFrame
Columns: [LYLTY_CARD_NBR, LIFESTAGE, PREMIUM_CUSTOMER]
Index: []
```


```
[40]: df_purchase_behaviour[['LIFESTAGE', 'PREMIUM_CUSTOMER']].describe()
```

```
[40]:      LIFESTAGE  PREMIUM_CUSTOMER
count      72637           72637
unique         7             3
top    RETIREES      Mainstream
freq      14805           29245
```

```
[41]: df_purchase_behaviour['LIFESTAGE'].unique()
```

```
[41]: array(['YOUNG SINGLES/COUPLES', 'YOUNG FAMILIES', 'OLDER SINGLES/COUPLES',
        'MIDAGE SINGLES/COUPLES', 'NEW FAMILIES', 'OLDER FAMILIES',
        'RETIREES'], dtype=object)
```


```
[42]: fig, ax = plt.subplots()
df_purchase_behaviour['LIFESTAGE'].value_counts().plot(ax=ax, kind='bar')
plt.xticks(rotation=45, ha='right')
plt.show()
```



```
[43]: df_purchase_behaviour['PREMIUM_CUSTOMER'].unique()
```

```
[43]: array(['Premium', 'Mainstream', 'Budget'], dtype=object)
```

```
[44]: fig, ax = plt.subplots()
df_purchase_behaviour['PREMIUM_CUSTOMER'].value_counts().plot(ax=ax, kind='bar')
plt.xticks(rotation=45, ha='right')
plt.show()
```



A few things we learned with this dataset: - There are no nulls or duplicates - Most of our client fall under the Mainstream category in 'PREMIUM_CUSTOMER' - Looking at the life stage of our clients, most of them are Retirees and New Families are the minority group.

1.4 Merge data

Now let's join both datasets so we can make a better analysis. Both datasets have the 'LYLTY_CARD_NBR' column, we have to do a left join on that.

```
[45]: joined_df = pd.
      ↪merge(df_transaction_data,df_purchase_behaviour,on='LYLTY_CARD_NBR',how='left')
      joined_df
```

```
[45]:
```

	index	OLD_DATE	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	\
0	0	43390	2018-10-17	1	1000	1	
1	1	43599	2019-05-14	1	1307	348	
2	2	43605	2019-05-20	1	1343	383	
3	3	43329	2018-08-17	2	2373	974	
4	4	43330	2018-08-18	2	2426	1038	
...	
246734	246736	43533	2019-03-09	272	272319	270088	
246735	246737	43325	2018-08-13	272	272358	270154	

246736	246738	43410	2018-11-06	272	272379	270187
246737	246739	43461	2018-12-27	272	272379	270188
246738	246740	43365	2018-09-22	272	272380	270189

	PROD_NBR		OLD_PROD_NAME	\
0	5	Natural Chip	Compny SeaSalt	175g
1	66		CCs Nacho Cheese	175g
2	61	Smiths Crinkle Cut	Chips Chicken	170g
3	69	Smiths Chip Thinly	S/Cream&Onion	175g
4	108	Kettle Tortilla ChpsHny&Jlpno	Chili	150g
...	...			
246734	89	Kettle Sweet Chilli And Sour Cream		175g
246735	74		Tostitos Splash Of Lime	175g
246736	51		Doritos Mexicana	170g
246737	42	Doritos Corn Chip Mexican	Jalapeno	150g
246738	74		Tostitos Splash Of Lime	175g

	PROD_NAME	BRND_NAME	PCK_SIZE	PROD_QTY	\
0	Chip	Compny SeaSalt	Natural	175	2
1		Nacho Cheese	CCs	175	3
2	Crinkle Cut	Chips Chicken	Smiths	170	2
3	Chip Thinly	S/Cream&Onion	Smiths	175	5
4	Tortilla ChpsHny&Jlpno	Chili	Kettle	150	3
...
246734	Sweet Chilli And Sour Cream		Kettle	175	2
246735		Splash Of Lime	Tostitos	175	1
246736		Mexicana	Doritos	170	2
246737	Corn Chip Mexican	Jalapeno	Doritos	150	2
246738		Splash Of Lime	Tostitos	175	2

	TOT_SALES		LIFESTAGE	PREMIUM_CUSTOMER
0	6.0	YOUNG SINGLES/COUPLES		Premium
1	6.3	MIDAGE SINGLES/COUPLES		Budget
2	2.9	MIDAGE SINGLES/COUPLES		Budget
3	15.0	MIDAGE SINGLES/COUPLES		Budget
4	13.8	MIDAGE SINGLES/COUPLES		Budget
...
246734	10.8	YOUNG SINGLES/COUPLES		Premium
246735	4.4	YOUNG SINGLES/COUPLES		Premium
246736	8.8	YOUNG SINGLES/COUPLES		Premium
246737	7.8	YOUNG SINGLES/COUPLES		Premium
246738	8.8	YOUNG SINGLES/COUPLES		Premium

[246739 rows x 15 columns]

Now that we have the joined data, let's make sure that there are no transactions without customer data, or if there is any customer with no transactions that we might have left out.

```
[46]: joined_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 246739 entries, 0 to 246738
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   index                 246739 non-null  int64  
 1   OLD_DATE              246739 non-null  int64  
 2   DATE                  246739 non-null  datetime64[ns]
 3   STORE_NBR             246739 non-null  int64  
 4   LYLTY_CARD_NBR        246739 non-null  int64  
 5   TXN_ID                246739 non-null  int64  
 6   PROD_NBR              246739 non-null  int64  
 7   OLD_PROD_NAME         246739 non-null  object  
 8   PROD_NAME             246739 non-null  object  
 9   BRND_NAME             246739 non-null  object  
10   PCK_SIZE              246739 non-null  int64  
11   PROD_QTY              246739 non-null  int64  
12   TOT_SALES             246739 non-null  float64 
13   LIFESTAGE             246739 non-null  object  
14   PREMIUM_CUSTOMER      246739 non-null  object  
dtypes: datetime64[ns](1), float64(1), int64(8), object(5)
memory usage: 30.1+ MB
```

```
[47]: df_purchase_behaviour[~df_purchase_behaviour.LYLTY_CARD_NBR.isin(joined_df.
↳ LYLTY_CARD_NBR.unique())]
```

```
[47]:
```

	LYLTY_CARD_NBR	LIFESTAGE	PREMIUM_CUSTOMER
21	1028	YOUNG SINGLES/COUPLES	Budget
78	1117	OLDER SINGLES/COUPLES	Mainstream
90	1137	MIDAGE SINGLES/COUPLES	Premium
95	1143	OLDER FAMILIES	Budget
100	1152	RETIREEES	Budget
...
72437	272164	YOUNG FAMILIES	Mainstream
72515	272276	YOUNG FAMILIES	Budget
72530	272295	RETIREEES	Mainstream
72547	272321	OLDER SINGLES/COUPLES	Premium
72608	880551	OLDER SINGLES/COUPLES	Premium

```
[1350 rows x 3 columns]
```

All transactions have customer behaviour data, but we don't have transactions for every customer. This might be due to the fact that we have data in the range of a year, and those customer might not have made any purchase in that time.

Let's save the joined dataset in a separate csv file for later use.

```
[48]: joined_df.to_csv('./data/QVI_transaction_and_customer_data.csv', index=False)
```

Now that data exploration is complete, we can move over to the analysis part.

1.5 Data analysis on customer segments

We'll start of by defining some metrics of interest to the client that we will try to come up with an answer: - How many customers are in each segment? What segment of clients represent most of our sales? - Chip brands preference and the favourite per segment. - Who spends the most on chips (total sales)? - How many chips are bought per customer by segment? - What's the average chip price by customer segment? - The customer's total spend over the period and total spend for each transaction to understand what proportion of their grocery spend is on chips - Proportion of customers in each customer segment overall to compare against the mix of customers who purchase chips

Let's start with sales by customers segments

```
[49]: prod_sales_sum = joined_df[['LIFESTAGE', 'PREMIUM_CUSTOMER', 'LYLTY_CARD_NBR']].
      →groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).count().
      →sort_values(by='LYLTY_CARD_NBR', ascending=False)
prod_sales_sum = prod_sales_sum.rename(columns={'LYLTY_CARD_NBR': 'N_CUSTOMERS'})
total_sales = prod_sales_sum.N_CUSTOMERS.sum()
prod_sales_sum['PROPORTION'] = prod_sales_sum.N_CUSTOMERS / total_sales * 100
prod_sales_sum
```

```
[49]:
```

LIFESTAGE	PREMIUM_CUSTOMER	N_CUSTOMERS	PROPORTION
OLDER FAMILIES	Budget	21514	8.719335
RETIREEES	Mainstream	19970	8.093573
YOUNG SINGLES/COUPLES	Mainstream	19544	7.920920
YOUNG FAMILIES	Budget	17763	7.199105
OLDER SINGLES/COUPLES	Budget	17172	6.959581
	Mainstream	17061	6.914594
	Premium	16559	6.711140
RETIREEES	Budget	14225	5.765201
OLDER FAMILIES	Mainstream	13241	5.366399
RETIREEES	Premium	12236	4.959086
YOUNG FAMILIES	Mainstream	11947	4.841959
MIDAGE SINGLES/COUPLES	Mainstream	11095	4.496654
YOUNG FAMILIES	Premium	10784	4.370610
OLDER FAMILIES	Premium	10403	4.216196
YOUNG SINGLES/COUPLES	Budget	8573	3.474522
MIDAGE SINGLES/COUPLES	Premium	7612	3.085041
YOUNG SINGLES/COUPLES	Premium	5852	2.371737
MIDAGE SINGLES/COUPLES	Budget	4691	1.901199
NEW FAMILIES	Budget	2824	1.144529
	Mainstream	2185	0.885551


Premium

1488

0.603066

```
[50]: prod_sales_sum['N_CUSTOMERS'].plot.bar()
```

```
[50]: <AxesSubplot: xlabel='LIFESTAGE, PREMIUM_CUSTOMER'>
```



```
[51]: prod_sales_sum = joined_df[['LIFESTAGE', 'PREMIUM_CUSTOMER', 'TOT_SALES']].  
      ↳groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).sum().sort_values(by='TOT_SALES',  
      ↳ascending=False)  
total_sales = prod_sales_sum.TOT_SALES.sum()
```


```
prod_sales_sum['PROPORTION_SALES'] = prod_sales_sum.TOT_SALES / total_sales * 100
prod_sales_sum
```

```
[51]:
```

		TOT_SALES	PROPORTION_SALES
LIFESTAGE	PREMIUM_CUSTOMER		
OLDER FAMILIES	Budget	156863.75	8.689686
YOUNG SINGLES/COUPLES	Mainstream	147582.20	8.175521
RETIREEES	Mainstream	145168.95	8.041836
YOUNG FAMILIES	Budget	129717.95	7.185906
OLDER SINGLES/COUPLES	Budget	127833.60	7.081520
	Mainstream	124648.50	6.905077
	Premium	123531.55	6.843202
RETIREEES	Budget	105916.30	5.867381
OLDER FAMILIES	Mainstream	96413.55	5.340963
RETIREEES	Premium	91296.65	5.057505
YOUNG FAMILIES	Mainstream	86338.25	4.782828
MIDAGE SINGLES/COUPLES	Mainstream	84734.25	4.693972
YOUNG FAMILIES	Premium	78571.70	4.352589
OLDER FAMILIES	Premium	75242.60	4.168169
YOUNG SINGLES/COUPLES	Budget	57122.10	3.164358
MIDAGE SINGLES/COUPLES	Premium	54443.85	3.015993
YOUNG SINGLES/COUPLES	Premium	39052.30	2.163357
MIDAGE SINGLES/COUPLES	Budget	33345.70	1.847231
NEW FAMILIES	Budget	20607.45	1.141578
	Mainstream	15979.70	0.885218
	Premium	10760.80	0.596110

```
[52]: prod_sales_sum['TOT_SALES'].plot.bar()
```

```
[52]: <AxesSubplot:xlabel='LIFESTAGE,PREMIUM_CUSTOMER'>
```



With this we now know that **our top segment** are **OLDER FAMILIES/Budget**, **RETIRES/Mainstream** and **YOUNG SINGLES/COUPLES/Mainstream**, this contributes to there being more sales to these customer segments.


```
[53]: joined_df[['LIFESTAGE', 'PREMIUM_CUSTOMER', 'BRND_NAME']].
      ↳groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).describe()
```

```
[53]:
```

	BRND_NAME			
	count	unique	top	freq
LIFESTAGE				
PREMIUM_CUSTOMER				

MIDAGE SINGLES/COUPLES	Budget	4691	26	Kettle	713
	Mainstream	11095	26	Kettle	2136
	Premium	7612	26	Kettle	1206
NEW FAMILIES	Budget	2824	26	Kettle	510
	Mainstream	2185	26	Kettle	414
	Premium	1488	26	Kettle	247
OLDER FAMILIES	Budget	21514	26	Kettle	3320
	Mainstream	13241	26	Kettle	2019
	Premium	10403	26	Kettle	1512
OLDER SINGLES/COUPLES	Budget	17172	26	Kettle	3065
	Mainstream	17061	26	Kettle	2835
	Premium	16559	26	Kettle	2947
RETIREEES	Budget	14225	26	Kettle	2592
	Mainstream	19970	26	Kettle	3386
	Premium	12236	26	Kettle	2216
YOUNG FAMILIES	Budget	17763	26	Kettle	2743
	Mainstream	11947	26	Kettle	1789
	Premium	10784	26	Kettle	1745
YOUNG SINGLES/COUPLES	Budget	8573	26	Kettle	1211
	Mainstream	19544	26	Kettle	3844
	Premium	5852	26	Kettle	838

```
[54]: fig, ax = plt.subplots(figsize=(15,5))
joined_df['BRND_NAME'].value_counts().plot(ax=ax, kind='bar')
plt.xticks(rotation=45, ha='right')
plt.show()
```



```
[55]: joined_df[['PREMIUM_CUSTOMER', 'BRND_NAME']].groupby(['PREMIUM_CUSTOMER']).
      ↳describe()
```

```
[55]:
```

	BRND_NAME	count	unique	top	freq

PREMIUM_CUSTOMER				
Budget	86762	26	Kettle	14154
Mainstream	95043	26	Kettle	16423
Premium	64934	26	Kettle	10711

With the table and graph we see that **Kettle is the top brand preferred by most groups and is top saled brand overall**, even accross those with budget and premium purchase behaviour.

```
[56]: prod_sales_sum = joined_df[['LYLTY_CARD_NBR', 'TOT_SALES']].
      ↪groupby('LYLTY_CARD_NBR').sum().sort_values(by='TOT_SALES', ascending=False).
      ↪head(10)
      pd.merge(prod_sales_sum, df_purchase_behaviour, on='LYLTY_CARD_NBR', how='left')
```

```
[56]:  LYLTY_CARD_NBR  TOT_SALES  LIFESTAGE PREMIUM_CUSTOMER
0         230078      138.6    OLDER FAMILIES      Budget
1         58361      124.8    YOUNG FAMILIES      Budget
2         63197      122.6    OLDER FAMILIES      Budget
3        162039      121.6    OLDER FAMILIES    Mainstream
4        179228      120.8    YOUNG FAMILIES      Budget
5        199157      118.8    YOUNG FAMILIES      Premium
6          3153      116.4  MIDAGE SINGLES/COUPLES      Premium
7         95048      115.1  YOUNG SINGLES/COUPLES    Mainstream
8          5168      114.8    OLDER FAMILIES    Mainstream
9         23192      114.7    OLDER FAMILIES      Budget
```

Here there is a top 10 list of our customer based on their spendings. Interestingly, **our top costumer doesn't have a premium but a butget purchasing behaviour**

```
[57]: joined_df[['LIFESTAGE', 'PREMIUM_CUSTOMER', 'PROD_QTY']].
      ↪groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).mean().sort_values(by='PROD_QTY',
      ↪ascending=False)
```

```
[57]:  LIFESTAGE  PREMIUM_CUSTOMER  PROD_QTY
OLDER FAMILIES  Mainstream      1.948795
               Premium         1.945496
               Budget         1.945384
YOUNG FAMILIES  Mainstream      1.941408
               Budget         1.941226
               Premium         1.938149
OLDER SINGLES/COUPLES  Budget      1.914920
                   Premium      1.913944
MIDAGE SINGLES/COUPLES  Mainstream      1.911942
OLDER SINGLES/COUPLES  Mainstream      1.911201
RETIREES               Premium      1.901438
MIDAGE SINGLES/COUPLES  Budget         1.893626
RETIREES               Budget         1.893286
MIDAGE SINGLES/COUPLES  Premium         1.891750
```

RETIREEES	Mainstream	1.886680
NEW FAMILIES	Premium	1.860887
	Mainstream	1.858124
	Budget	1.855878
YOUNG SINGLES/COUPLES	Mainstream	1.853510
	Budget	1.808002
	Premium	1.807075

Older and young families in general buy more chips per customer. This helps they beign on top of our sales charts from before.

```
[58]: joined_df['AVG_PRICE'] = joined_df.TOT_SALES / joined_df.PROD_QTY
      joined_df[['LIFESTAGE', 'PREMIUM_CUSTOMER', 'AVG_PRICE']].
      ↪groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER']).mean().sort_values(by='AVG_PRICE',
      ↪ascending=False)
```

```
[58]:
```

		AVG_PRICE
LIFESTAGE	PREMIUM_CUSTOMER	
YOUNG SINGLES/COUPLES	Mainstream	4.065642
MIDAGE SINGLES/COUPLES	Mainstream	3.994241
RETIREEES	Budget	3.924404
	Premium	3.920942
NEW FAMILIES	Budget	3.917688
	Mainstream	3.916133
OLDER SINGLES/COUPLES	Premium	3.893236
	Budget	3.882096
NEW FAMILIES	Premium	3.872110
RETIREEES	Mainstream	3.844294
OLDER SINGLES/COUPLES	Mainstream	3.814665
MIDAGE SINGLES/COUPLES	Premium	3.770698
YOUNG FAMILIES	Premium	3.762150
	Budget	3.760737
OLDER FAMILIES	Budget	3.745340
MIDAGE SINGLES/COUPLES	Budget	3.743328
OLDER FAMILIES	Mainstream	3.737077
YOUNG FAMILIES	Mainstream	3.724533
OLDER FAMILIES	Premium	3.717000
YOUNG SINGLES/COUPLES	Premium	3.665414
	Budget	3.657366

In the table above we can se the average chip price by customer segment. Our Young and Midage Single/Couples Mainstream tend to pay more for chips than their budget and premium counter-parts.

As the difference in average price per unit isn't large, we can check if this difference is statistically different.

```
[59]: from scipy.stats import ttest_ind

filter_df = joined_df[(joined_df['LIFESTAGE'] == 'YOUNG SINGLES/COUPLES') |
↳ (joined_df['LIFESTAGE'] == 'MIDAGE SINGLES/COUPLES')]
mainstream_group = filter_df[filter_df.PREMIUM_CUSTOMER == 'Mainstream']
budget_premium_group = filter_df[~(filter_df.PREMIUM_CUSTOMER == 'Mainstream')]

t2, p2 = ttest_ind(mainstream_group.AVG_PRICE, budget_premium_group.AVG_PRICE)

print('P-Value = {}'.format(p2))
```

P-Value = 2.235645611549355e-309

The t-test results in a p-value close to zero, i.e. the unit price for mainstream, young and mid-age singles and couples ARE significantly higher than that of budget or premium, young and midage singles and couples.

With all that we have seen so far we should focus on retaining those customers that tend to pay more per chips, **that beeing Young and Midage Single/Couples Mainstream**. Also having in mind that their preferd brand is **Kettle**.

```
[60]: filter_df.PCK_SIZE = filter_df.PCK_SIZE.astype('str')
filter_df.groupby('LIFESTAGE')['PCK_SIZE'].value_counts()
```

```
[60]: LIFESTAGE      PCK_SIZE
MIDAGE SINGLES/COUPLES 175      6334
                        150      3755
                        134      2389
                        110      2223
                        170      1835
                        165      1448
                        330      1220
                        270       649
                        380       628
                        210       613
                        200       388
                        135       309
                        160       290
                        250       282
                        190       266
                        90       249
                        220       152
                        70       141
                        180       119
                        125       108
YOUNG SINGLES/COUPLES 175     8953
                        150     5403
                        134     3684
```

110	3227
170	2751
165	2077
330	1803
380	955
270	947
210	913
200	538
135	452
250	448
160	388
190	369
90	361
70	182
220	178
125	177
180	163

Name: PCK_SIZE, dtype: int64

We should also have in mind that our segment of interest prefers medium size packs of chips, **175g** is preferred for both groups.